

CS446 Introduction to Machine Learning (Spring 2015)  
University of Illinois at Urbana-Champaign  
<http://courses.engr.illinois.edu/cs446>

# LECTURE 2:

## SUPERVISED LEARNING

Prof. Julia Hockenmaier  
juliaahr@illinois.edu

# Class admin

Are you on Piazza?

Is everybody registered for the class?

HWO is out (not graded)

<http://courses.engr.illinois.edu/cs446/Homework/HWO/HWO.pdf>

Email alias for CS446 staff:

[cs446-staff@mx.uillinois.edu](mailto:cs446-staff@mx.uillinois.edu)

# Learning scenarios

## Supervised learning:

The focus of CS446

Learning to predict labels from correctly labeled data

## Unsupervised learning:

Learning to find hidden structure (e.g. clusters) in input data

## Semi-supervised learning:

Learning to predict labels from (a little) labeled and (a lot of) unlabeled data

## Reinforcement learning:

Learning to act through feedback for actions (rewards/punishments) from the environment

# The Badges game

+ Naoki Abe

- Eric Baum

Conference attendees to the 1994 Machine Learning conference were given  
name badges labeled with + or −.

What function was used to assign these labels?

# The supervised learning task

Given a **labeled training data** set  
of  $N$  items  $\mathbf{x}_n \in \mathcal{X}$  with labels  $y_n \in \mathcal{Y}$

$$\mathcal{D}^{\text{train}} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)\}$$

( $y_n$  is determined by some unknown target function  $f(\mathbf{x})$ )

Return **a model**  $g: \mathcal{X} \mapsto \mathcal{Y}$  that is a good  
approximation of  $f(\mathbf{x})$

( $g$  should assign correct labels  $y$  to **unseen**  $\mathbf{x} \notin \mathcal{D}^{\text{train}}$ )

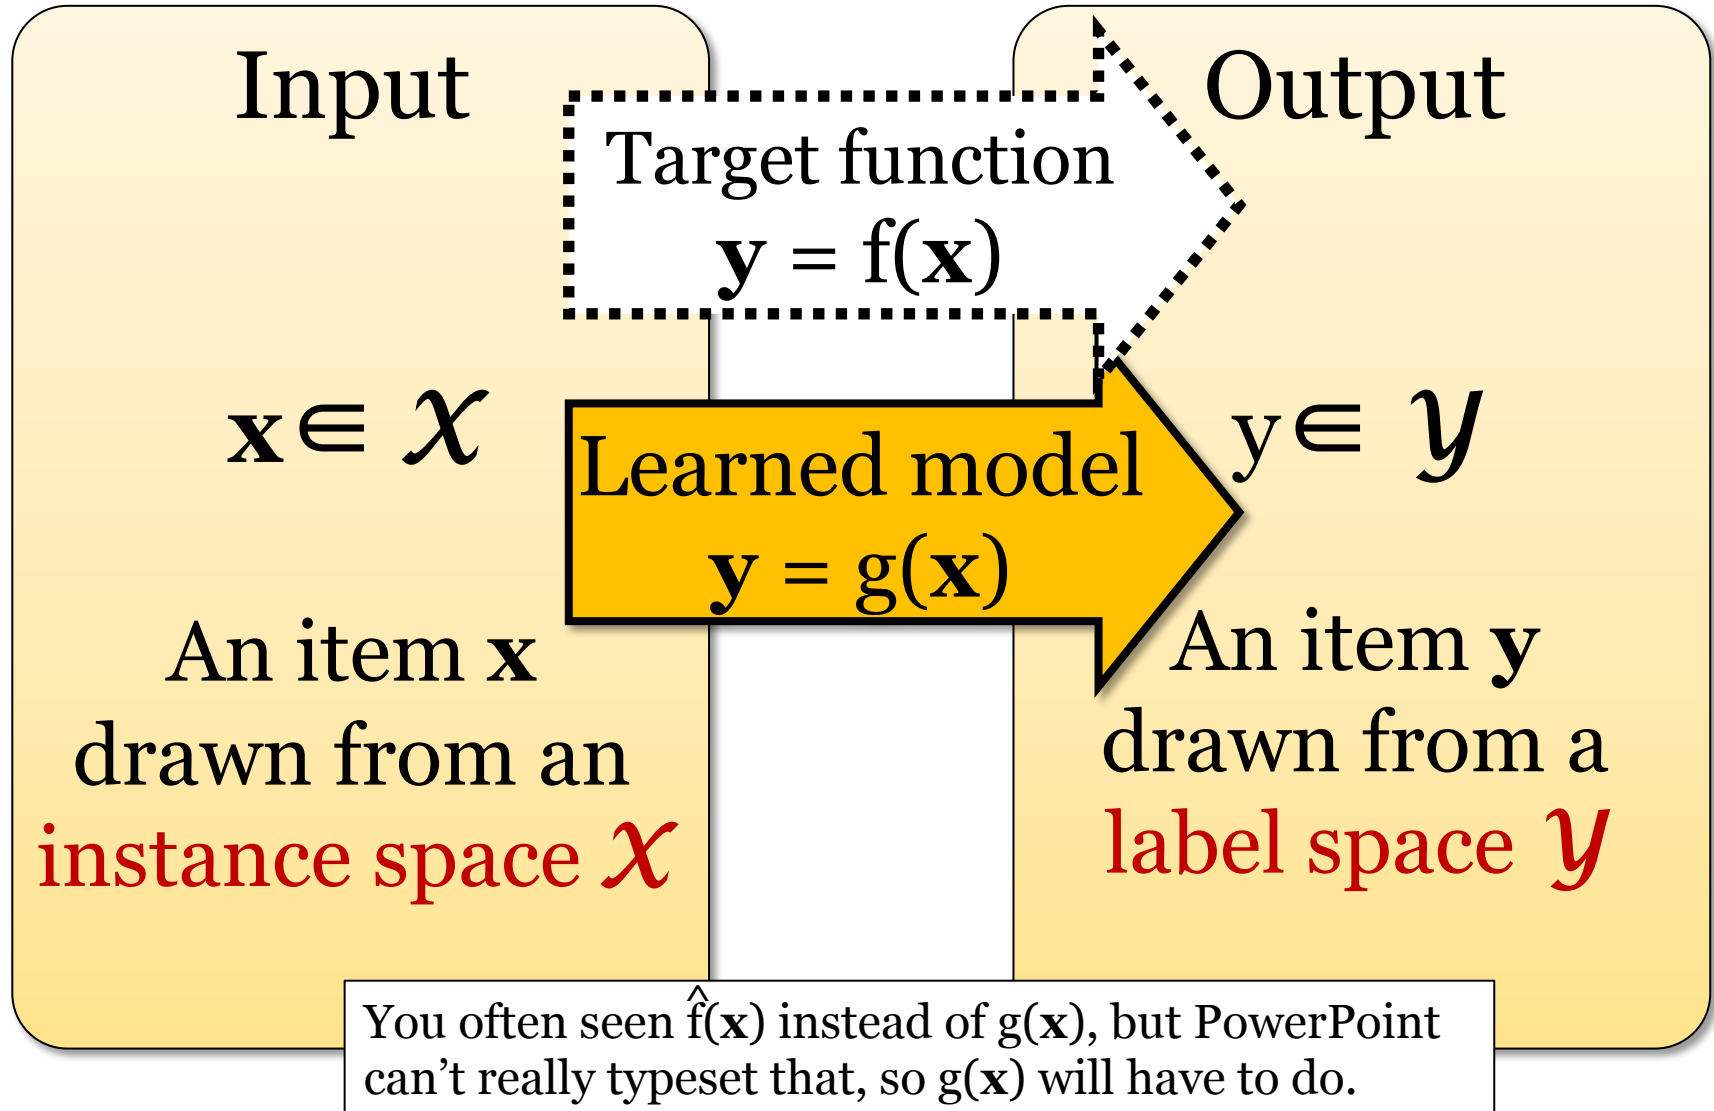
# Supervised learning terms

Input items/data points  $\mathbf{x}_n \in \mathcal{X}$  (e.g. emails) are drawn from an **instance space  $\mathcal{X}$**


Output labels  $y_n \in \mathcal{Y}$  (e.g. 'spam'/'nospam') are drawn from a **label space  $\mathcal{Y}$**

Every data point  $\mathbf{x}_n \in \mathcal{X}$  has a *single* correct label  $y_n \in \mathcal{Y}$ , defined by an (unknown) **target function  $f(\mathbf{x}) = y$**

# Supervised learning



# Supervised learning: Training



Give the learner examples in  $\mathcal{D}^{\text{train}}$

The learner returns a model  $g(\mathbf{x})$



# Training data

+ Naoki Abe	+ Peter Bartlett	+ Carla E. Brodley
- Myriam Abramson	- Eric Baum	+ Nader Bshouty
+ David W. Aha	+ Welton Becket	- Wray Buntine
+ Kamal M. Ali	- Shai Ben-David	- Andrey Burago
- Eric Allender	+ George Berg	+ Tom Bylander
+ Dana Angluin	+ Neil Berkman	+ Bill Byrne
- Chidanand Apte	+ Malini Bhandaru	- Claire Cardie
+ Minoru Asada	+ Bir Bhanu	+ John Case
+ Lars Asker	+ Reinhard Blasig	+ Jason Catlett
+ Javed Aslam	- Avrim Blum	- Philip Chan
+ Jose L. Balcazar	- Anselm Blumer	- Zhixiang Chen
- Cristina Baroglio	+ Justin Boyan	- Chris Darken

# Supervised learning: Testing

Labeled  
Test Data

$\mathcal{D}^{\text{test}}$

$(\mathbf{x}'_1, y'_1)$

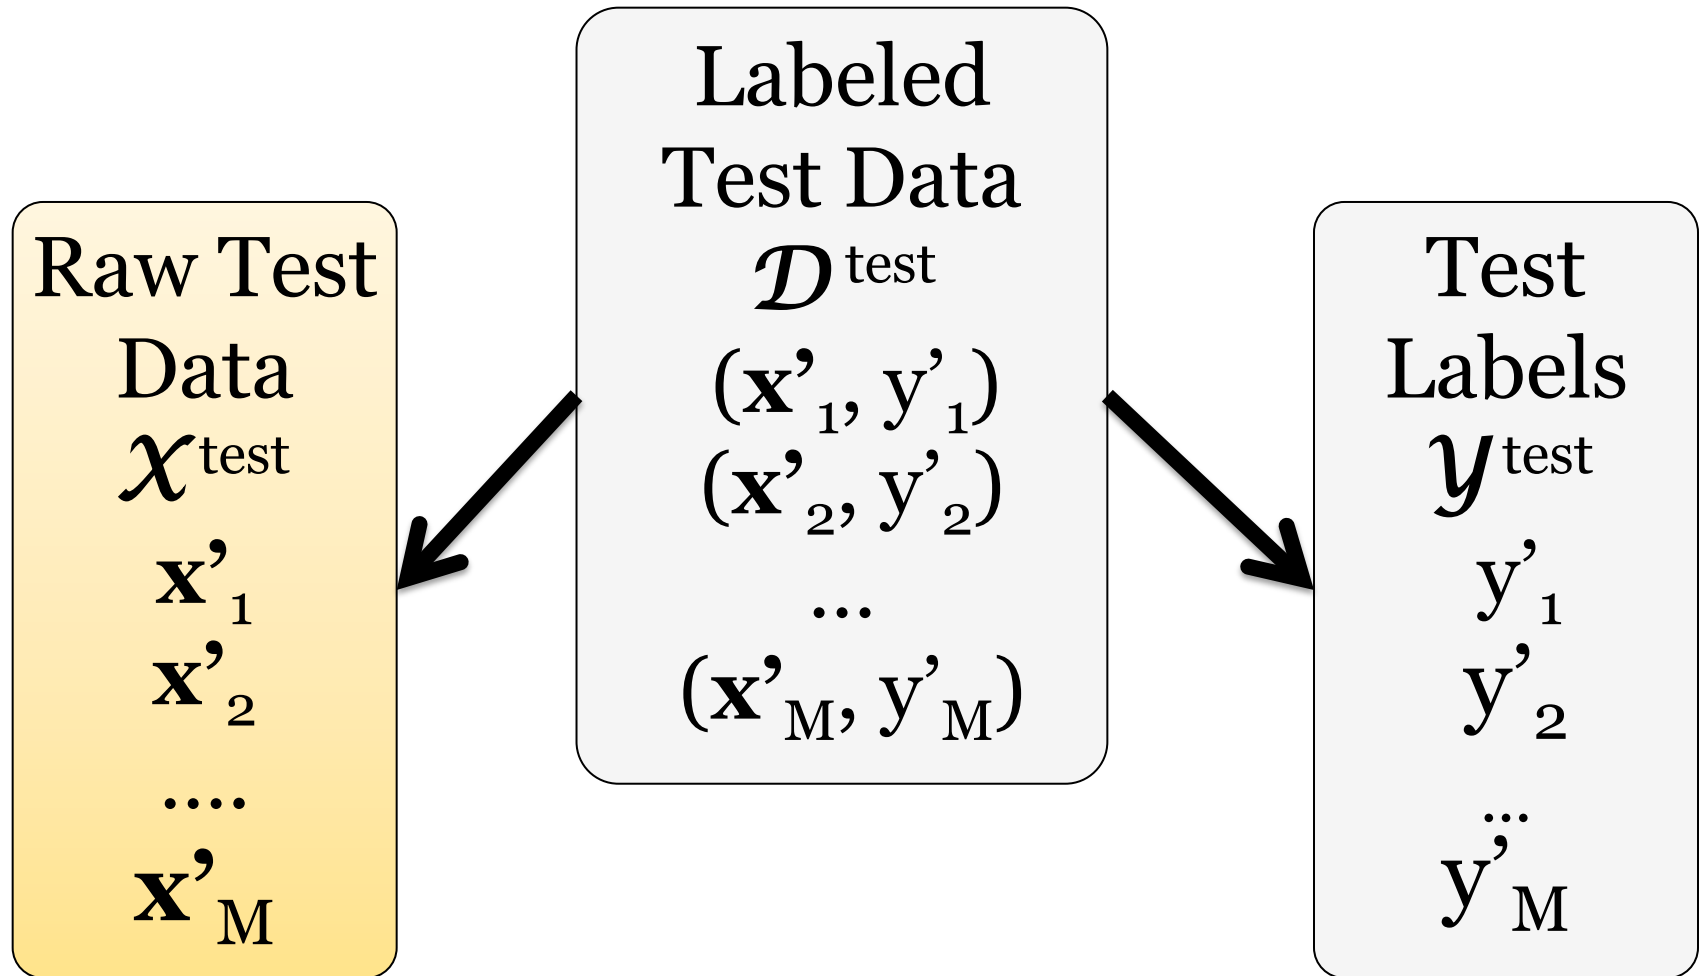
$(\mathbf{x}'_2, y'_2)$

...

$(\mathbf{x}'_M, y'_M)$

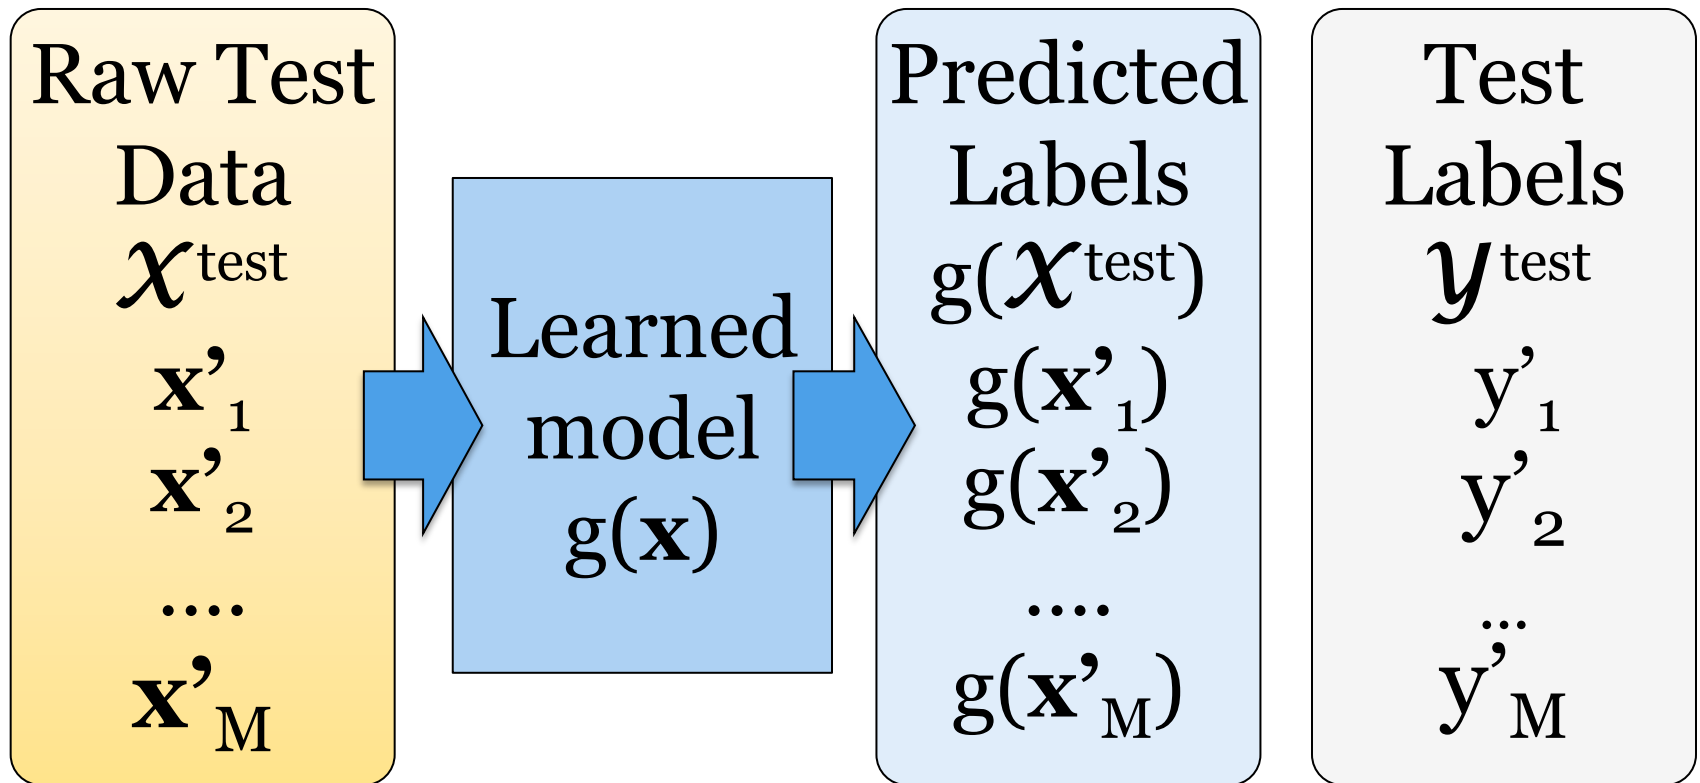
Reserve some labeled data for testing

# Supervised learning: Testing



# Supervised learning: Testing

Apply the model to the raw test data




# Raw test data

Gerald F. DeJong  
Chris Drummond  
Yolanda Gil  
Attilio Giordana  
Jiarong Hong

J. R. Quinlan  
Priscilla Rasmussen  
Dan Roth  
Yoram Singer  
Lyle H. Ungar

# Supervised learning: Testing

Evaluate the model by comparing the predicted labels against the test labels



# Labeled test data

- + Gerald F. DeJong
- Chris Drummond
- + Yolanda Gil
- Attilio Giordana
- + Jiarong Hong
- J. R. Quinlan
- Priscilla Rasmussen
- + Dan Roth
- + Yoram Singer
- Lyle H. Ungar

# Evaluating supervised learners

Use a **test data set** that is disjoint from  $\mathcal{D}^{\text{train}}$

$$\mathcal{D}^{\text{test}} = \{(\mathbf{x}'_1, y'_1), \dots, (\mathbf{x}'_M, y'_M)\}$$

The learner has **not seen the test items** during learning.

Split your labeled data into two parts: test and training.

Take all items  $\mathbf{x}'_i$  in  $\mathcal{D}^{\text{test}}$  and compare the **predicted  $f(\mathbf{x}'_i)$**  with **the correct  $y'_i$** .

This requires an **evaluation metric** (e.g. accuracy).



# Using supervised learning

- What is our **instance space**?

Gloss: What kind of features are we using?

- What is our **label space**?

Gloss: What kind of learning task are we dealing with?

- What is our **hypothesis space**?

Gloss: What kind of model are we learning?

- What **learning algorithm** do we use?


Gloss: How do we learn the model from the labeled data?

- (What is our **loss function**/evaluation metric?)

Gloss: How do we measure success?

# 1. The instance space

# 1. The instance space $\mathcal{X}$



Designing an appropriate instance space  $\mathcal{X}$  is crucial for how well we can predict  $\mathbf{y}$ .

# 1. The instance space $\mathcal{X}$

When we apply machine learning to a task, we first need to *define* the instance space  $\mathcal{X}$ .

Instances  $\mathbf{x} \in \mathcal{X}$  are defined by **features**:

- **Boolean features:**

Does this email contain the word ‘money’?

- **Numerical features:**

How often does ‘money’ occur in this email?

What is the width/height of this bounding box?

# What's $\mathcal{X}$ for the Badges game?

Possible features:

- Gender/age/country of the person?
- Length of their first or last name?
- Does the name contain letter 'x'?
- How many vowels does their name contain?
- Is the n-th letter a vowel?


# $\mathcal{X}$ as a vector space

$\mathcal{X}$  is an  $N$ -dimensional vector space (e.g.  $\mathbb{R}^N$ )

Each dimension = one feature.

Each  $\mathbf{x}$  is a **feature vector** (hence the boldface  $\mathbf{x}$ ).

Think of  $\mathbf{x} = [x_1 \dots x_N]$  as a point in  $\mathcal{X}$ :



# From feature templates to vectors

When designing features, we often think in terms of **templates**, not individual features:

**What is the 2nd letter?**

N **a** oki  $\rightarrow [1\ 0\ 0\ 0\ \dots]$

A **b** e  $\rightarrow [0\ 1\ 0\ 0\ \dots]$

S **c** rooge  $\rightarrow [0\ 0\ 1\ 0\ \dots]$

**What is the  $i$ -th letter?**

**A****b****e**  $\rightarrow [1\ 0\ 0\ 0\ 0\ \dots\ 0\ 1\ 0\ 0\ 0\ 0\ \dots\ 0\ 0\ 0\ 0\ 1\ \dots]$

# Good features are essential

The choice of features is crucial  
for how well a task can be learned.

In many application areas (language, vision, etc.),  
a lot of work goes into designing suitable features.  
This requires domain expertise.


CS446 can't teach you what specific features  
to use for your task.

But we will touch on some general principles



## 2. The label space

## 2. The label space $\mathcal{Y}$



The label space  $\mathcal{Y}$  determines *what kind of supervised learning task* we are dealing with

# Supervised learning tasks I

Output labels  $y \in \mathcal{Y}$  are **categorical**: The focus of CS446

- Binary classification: Two possible labels
- Multiclass classification:  $k$  possible labels

Output labels  $y \in \mathcal{Y}$  are **structured objects**  
(sequences of labels, parse trees, etc.)

- Structure learning (e.g. CS546)

# Supervised learning tasks II

Output labels  $y \in \mathcal{Y}$  are **numerical**:

- **Regression** (linear/polynomial):

Labels are continuous-valued

Learn a linear/polynomial function  $f(x)$

- **Ranking**:


Labels are ordinal

Learn an ordering  $f(x_1) > f(x_2)$  over input

# 3. Models

(The hypothesis space)

### 3. The model $g(\mathbf{x})$



We need to choose what *kind* of model we want to learn

# More terminology

For classification tasks ( $\mathcal{Y}$  is categorical, e.g.  $\{0, 1\}$ , or  $\{0, 1, \dots, k\}$ ), the model is called a **classifier**.

For **binary classification tasks** ( $\mathcal{Y} = \{0, 1\}$ ), we often think of the two values of  $\mathcal{Y}$  as **Boolean** (0 = false, 1 = true), and call the target function  $f(\mathbf{x})$  to be learned a **concept**

# A learning problem

	$x_1$	$x_2$	$x_3$	$x_4$	$y$
1	0	0	1	0	0
2	0	1	0	0	0
3	0	0	1	1	1
4	1	0	0	1	1
5	0	1	1	0	0
6	1	1	0	0	0
7	0	1	0	1	0



# A learning problem

Each  $\mathbf{x}$  has 4 bits:  $|\mathcal{X}| = 2^4 = 16$

Since  $\mathcal{Y} = \{0, 1\}$ , each  $f(\mathbf{x})$  defines one subset of  $\mathcal{X}$

$\mathcal{X}$  has  $2^{16} = 65536$  subsets:

There are  $2^{16}$  possible  $f(\mathbf{x})$   
( $2^9$  are consistent with our data)

We would need to see all of  $\mathcal{X}$  to learn  $f(\mathbf{x})$

# A learning problem

We would need to see all of  $\mathcal{X}$  to learn  $f(\mathbf{x})$

- Easy with  $|\mathcal{X}|=16$
- Not feasible in general  
(for any real-world problems)
- Learning = generalization,  
not memorization of the training data

# The hypothesis space $\mathcal{H}$

There are  $|\mathcal{Y}|^{|\mathcal{X}|}$  possible functions  $f(\mathbf{x})$  from the instance space  $\mathcal{X}$  to the label space  $\mathcal{Y}$ .

Learners typically consider *only a subset of the functions from  $\mathcal{X}$  to  $\mathcal{Y}$* .

This subset is called the *hypothesis space  $\mathcal{H}$* .

$$\mathcal{H} \subseteq |\mathcal{Y}|^{|\mathcal{X}|}$$

# Can we restrict $\mathcal{H}$ ?

	$x_1$	$x_2$	$x_3$	$x_4$	$y$
1	0	0	1	0	0
2	0	1	0	0	0
3	0	0	1	1	1
4	1	0	0	1	1
5	0	1	1	0	0
6	1	1	0	0	0
7	0	1	0	1	0

## Conjunctive clauses:

16 different conjunctions  
over  $\{x_1, x_2, x_3, x_4\}$ :

$$f(\mathbf{x}) = x_1$$

...

$$f(\mathbf{x}) = x_1 \wedge x_2 \wedge x_3 \wedge x_4$$

None is consistent with  
the data

# Can we restrict $\mathcal{H}$ ?

	$x_1$	$x_2$	$x_3$	$x_4$	$y$
1	0	0	1	0	0
2	0	1	0	0	0
3	0	0	1	1	1
4	1	0	0	1	1
5	0	1	1	0	0
6	1	1	0	0	0
7	0	1	0	1	0

## **n-of-m clauses:**

20 rules of the form

“ $y=1$  iff at least  $m$  of the following  $n$   $x_i$  are 1”

# Can we restrict $\mathcal{H}$ ?

	$x_1$	$x_2$	$x_3$	$x_4$	$y$
1	0	0	1	0	0
2	0	1	0	0	0
3	0	0	1	1	1
4	1	0	0	1	1
5	0	1	1	0	0
6	1	1	0	0	0
7	0	1	0	1	0

## **n-of-m clauses:**


20 rules of the form

“ $y=1$  iff at least  $m$  of the following  $n$   $x_i$  are 1”

## **Consistent hypothesis:**

“ $y=1$  if and only if at least 2 of  $\{x_1, x_3, x_4\}$  are 1”

# Classifiers in vector spaces



## Binary classification:

We assume  $f$  *separates* the positive and negative examples:

- Assign  $y = 1$  to all  $\mathbf{x}$  where  $f(\mathbf{x}) > 0$
- Assign  $y = 0$  to all  $\mathbf{x}$  where  $f(\mathbf{x}) < 0$

# Learning a classifier





## The learning task:

Find a function  $f(\mathbf{x})$  that best separates the (training) data

- What kind of function is  $f$ ?
- How do we define *best*?
- How do we find  $f$ ?



# Which model should we pick?



# Criteria for choosing models

## Accuracy:

Prefer models that make fewer mistakes


- We only have access to the training data
- But we care about accuracy on unseen (test) examples

## Simplicity (Occam's razor):

Prefer simpler models (e.g. fewer parameters).

- These (often) generalize better, and need less data for training.

# Linear classifiers




Many learning algorithms restrict the hypothesis space to **linear classifiers**:


$$f(\mathbf{x}) = w_0 + \mathbf{w}\mathbf{x}$$

# Linear Separability

Not all data sets are linearly separable:



Sometimes, feature transformations help:



# 4. The learning algorithm

# 4. The learning algorithm

The learning task:

Given a labeled training data set

$$\mathcal{D}^{\text{train}} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)\}$$

return a model (classifier)  $g: \mathcal{X} \mapsto \mathcal{Y}$

from the hypothesis space  $\mathcal{H} \subseteq |\mathcal{Y}|^{|\mathcal{X}|}$

The learning algorithm performs a **search in the hypothesis space  $\mathcal{H}$**  for the model  $g$ .

# Batch versus online training

## Batch learning:

The learner sees the complete training data, and only changes its hypothesis when it has seen **the entire training data set**.

## Online training:

The learner sees the training data one example at a time, and can change its hypothesis **with every new example**

# Practical issues



# How to use your labeled data

Split your data into two (or three) sets:

- Training data (often 70-90%)
- Test data (often 10-20%)
- Development data (10-20%)

You need to report performance on test data, but you are not allowed to look at it.

You are allowed to look at the development data (and use it to tweak parameters)

# Baseline performance

How difficult is your task?

You need to compare against a (reasonable) baseline (e.g. assign the majority class)

# Ablation studies

How important are the different features (feature templates) you have designed?


An ablation study compares models that use different subsets of the features/feature templates.

# Learning curves

How much training data do you need?

Has your model converged?

How does your performance change with the amount of training data?



Today's key  
concepts

# Lecture 2 key concepts

- Instance space (typically a vector space):  
each instance = one feature vector  $\mathbf{x} = (x_1 \dots x_n)$
- Hypothesis space (supervised learning):  
Subset of functions from instances to labels
- Linear classifiers:  
Only consider linear functions  $g(\mathbf{x}) = w_0 + \mathbf{w}\mathbf{x}$
- Learning algorithms:  
Online vs. batch learning
- Training vs. test vs. development data